



Tech Quicky - 2022 Jan

Introduction to NLP

Word Vectors

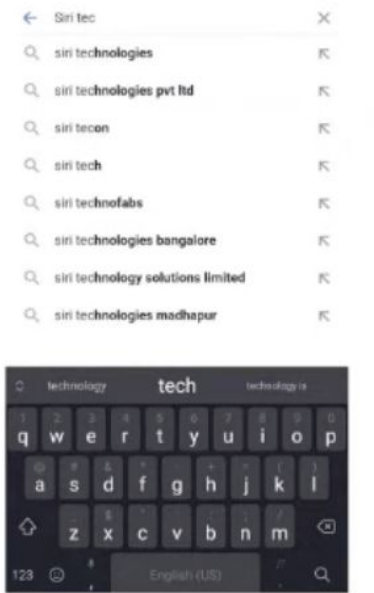


Natural language processing (NLP)


process/analyze natural language data (text/speech)
with algorithms and computer systems



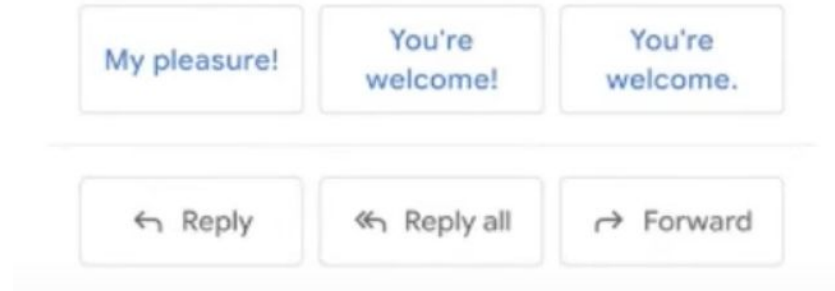
Examples



Autocomplete / Autocorrect



Machine Translation




Text / Response Generation



Huge and booming field

Only looking at a thin slice today



Word Vectors

Distributional Word Vectors

Word Embeddings

Neural Word Representation



Problem statement

How can we

- get computers to understand the language?
- represent word meaning?
- encode similarity?

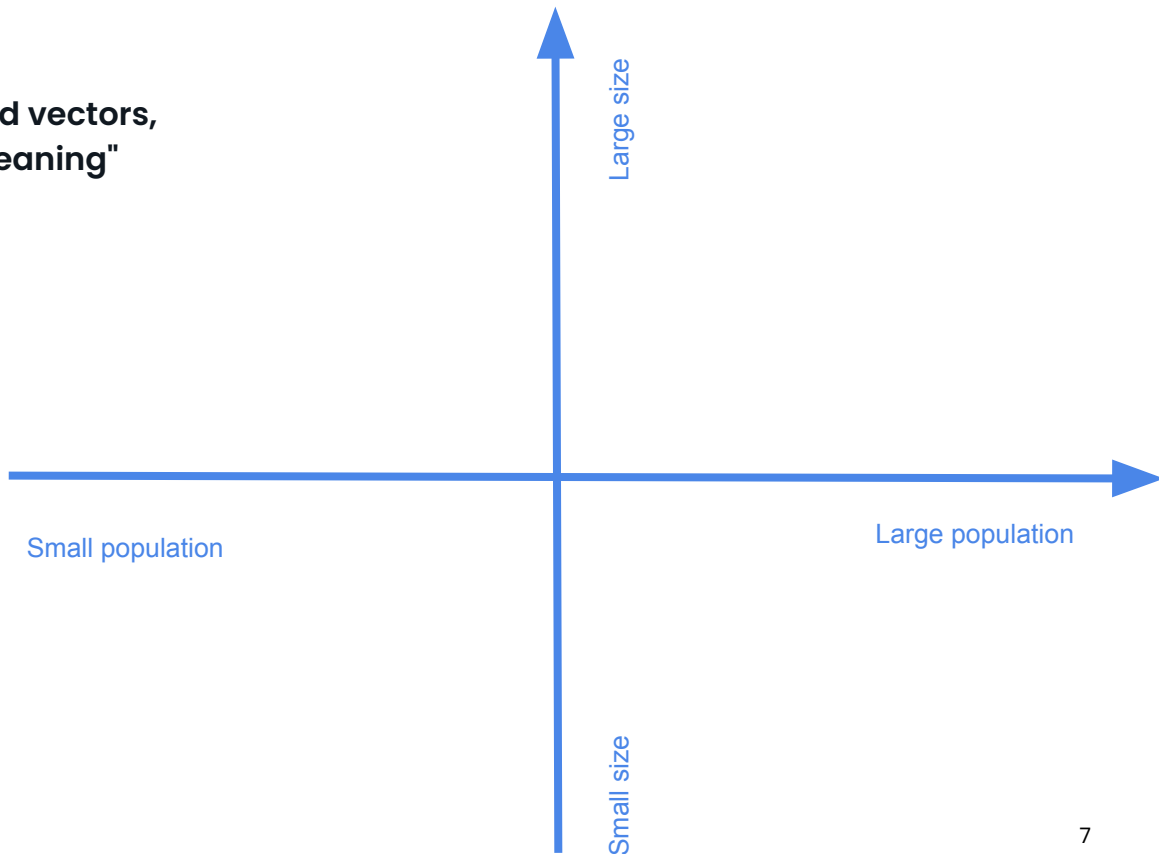


Word Vectors – intuition








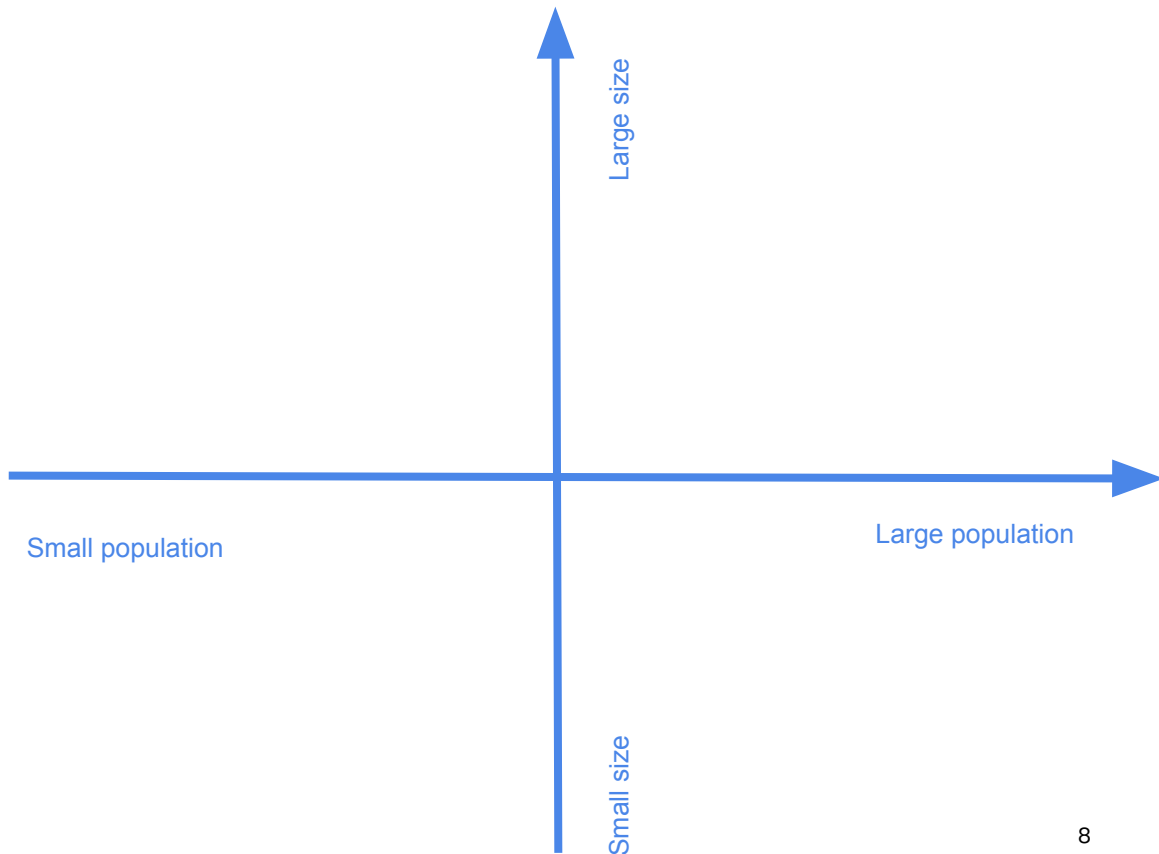
Word Vectors – Intuition

- let's represent words as fixed sized vectors,
- with dimensions along axis of "meaning"









Word Vectors – Intuition

Country	Population (mil pp)	Size (mil km ²)
 France	68	0.5
 Spain	47	0.5
 Brazil	214	8.5
 Iceland	0.4	0.1
 Canada	38	9



Word Vectors – Intuition

Country	Population (mil pp)	Size (mil km ²)
 France	68	0.5
 Spain	47	0.5
 Brazil	214	8.5
 Iceland	0.4	0.1
 Canada	38	9




Word Vectors – Intuition


Measured similarity by comparing values along axis


 [0.9, 0.8]

 [0.3, 0.2]

 [0.2, 0.2]

 [-.1, 0.9]

 [-.8, -.8]




Word Vectors – Intuition


Measured similarity by comparing values along axis

 [0.9, 0.8]

 [0.3, 0.2]

 [0.2, 0.2]

 [-.1, 0.9]

 [-.8, -.8]


Dot Product similarity

$$\text{dot}(a,b) = a \cdot b = x_1 \times x_2 + y_1 \times y_2$$

$$\text{France} \cdot \text{Spain} = 0.1$$

$$\text{Brazil} \cdot \text{Canada} = 0.63$$

$$\text{Brazil} \cdot \text{Norway} = -1.36$$



Word Vectors – Intuition

Normalize by size of vectors, measure similarity by angle

$$\text{🇧🇷} [0.9, 0.8] - \|\text{🇧🇷}\| = 1.2$$


$$\text{🇫🇷} [0.3, 0.2] - \|\text{🇫🇷}\| = 0.36$$

$$\text{🇪🇸} [0.2, 0.2] - \|\text{🇪🇸}\| = 0.28$$

$$\text{🇨🇦} [-.1, 0.9] - \|\text{🇨🇦}\| = 0.9$$

$$\text{🇩🇰} [-.8, -.8] - \|\text{🇩🇰}\| = 1.13$$

$\|\mathbf{a}\| = \sqrt{\mathbf{a} \cdot \mathbf{a}}$ (Euclidean length of the vector)



Word Vectors – Intuition

Normalize by size of vectors, measure similarity by angle

$$\text{Brazil} [0.9, 0.8] - \|\text{Brazil}\| = 1.2$$

$$\text{France} [0.3, 0.2] - \|\text{France}\| = 0.36$$

$$\text{Spain} [0.2, 0.2] - \|\text{Spain}\| = 0.28$$

$$\text{Canada} [-.1, 0.9] - \|\text{Canada}\| = 0.9$$

$$\text{Norway} [-.8, -.8] - \|\text{Norway}\| = 1.13$$

Cosine similarity – normalized dot product


$$\text{cosine}(a,b) = a \cdot b / (\|a\| \times \|b\|)$$

$$\text{cosine}(\text{France}, \text{Spain}) = 0.1 / 0.1008 = 0.99$$

$$\text{cosine}(\text{Brazil}, \text{Canada}) = 0.63 / 1.08 = 0.58$$

$$\text{cosine}(\text{Brazil}, \text{Norway}) = -1.36 / 1.36 = -0.99$$

Values go from $[-1, 1]$



Word Vectors – calculation

How could we calculate meaningful word vectors automatically?



Idea 1 – Distributional Semantics

words occurring in the same contexts
tend to have similar meanings

"a word is characterized by the company it keeps"
J.R. Firth in 1957

Let's define a word's meaning by

- its context
- its neighboring words



Example:

*To start my day, I drink hot **coffee** in the mornings.*

*Sipping some warm **tea** helps me wake up in the morning.*

Context:

coffee : drink, hot, morning, start my day

tea : sipping, warm, wake up, morning

coffee & tea


- both a liquid
- are consumed warm
- early in the day to activate people

Idea 2 – Self-supervised learning

We have lots of examples of language in use, available to us:

- Digitized books, newspapers
- Wikipedia, forums
- Social Media posts (reddit, twitter)

Let's train a neural network to calculate the word vectors
– using these existing bodies of text as labelled data



Steps to calculate word vector:

1. Start with random word vectors assigned to each word
2. Go through each position in the training data (text corpus)
3. Take a center word and surrounding context words
4. Use word vector similarity to calculate probability of co-occurrence
5. Keep adjusting word vectors until you maximize the probability based on actual co-occurrences in the corpus



Timeline

2000s

- the idea of calculating word embedding using neural networks (Bengio 2003)
- word embedding are very useful for NLP tasks (Collobert 2007)

Early 2010s

- Efficient ways to calculate embeddings
 - Word2Vec paper - (Mikolov et al. 2013)
 - GloVe paper - (Pennington et al. 2014)

GloVe:

Global Vectors for Word Representation



Word Vectors – Demo

aka the “fun part”



Word Vectors – Usefulness

What are word vectors good at?
What are word vectors bad at?



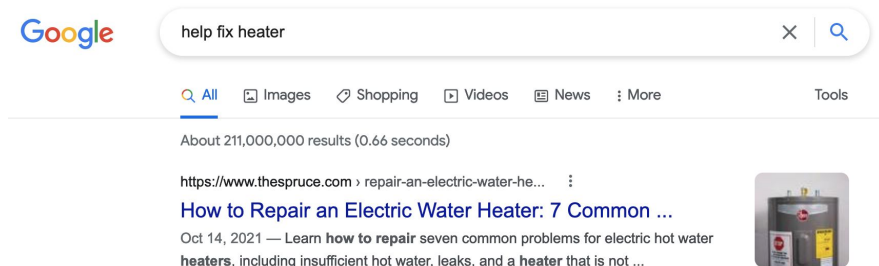
What are word vectors good at?

Automated dictionary generation (🇬🇧 → 🇪🇸)



Semantic search

- Classical search is lexical/keyword based vs
- Understanding query context, user intent, similarities



They revolutionized NLP

Almost all research in NLP start with word vectors and use it as a baseline to compare results to

In many NLP tasks state-of-the-art results are obtained using word vectors



What are word vectors bad at?

Static mappings

- one vector mapped to all occurrences of a word
- different meanings are mixed together
- vectors for the different meanings are averaged

Notebook

do we mean 📖 or 💻?

Park

are we talking about 🚗 or 🌳?

Rare words

- if a word is only present a few times in the corpus, the algorithms will have problems learning good representations
- rare words, have bad representation and frequently appear close to other, unrelated rare words

They have their weaknesses



What's next?

What comes after word vectors?



Dynamic, contextualized word vectors

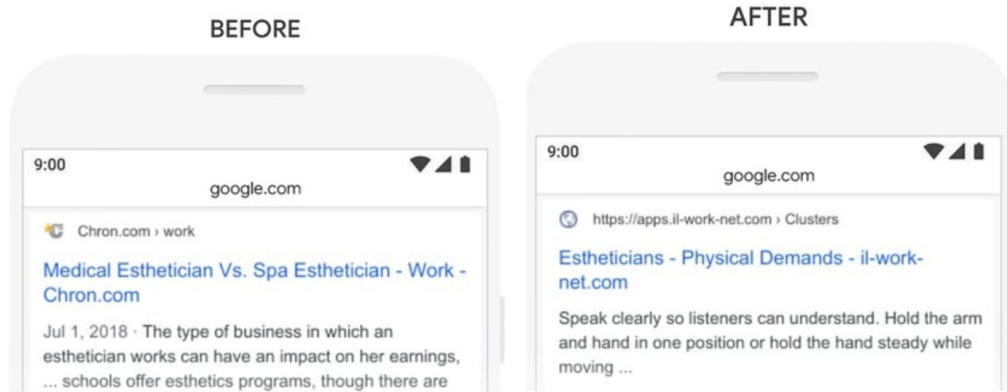
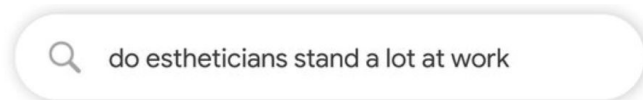
Compute a different vector for each occurrence of a word, based on its immediate context

Many improvements:

- Attention mechanism (Vaswani et al. 2017)
- Transformer-based models
 - Recurrent neural networks (RNN)
 - Sequence-to-Sequence(Seq2Seq)
 - Encoder-decoder networks

In 2019 [Google Search was updated with BERT](#) model resulting in a massive improvement:

“the biggest leap forward in the past five years, and one of the biggest leaps forward in the history of Search”



Take away

- A surprising result – word meaning can be represented well by large vectors of numbers
- These vectors can be calculated using a "simple" task of calculating and updating distributional similarities



Thank you!

Any questions?

